

# Deep Neural Network

Philip Hoddinott

December 17, 2018

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	The MNIST database . . . . .	2
1.2	Artificial neural network . . . . .	2
1.3	Neural Network Walkthrough . . . . .	3
1.3.1	Parameter Initialization . . . . .	3
1.3.2	Forward Propagation . . . . .	3
1.3.3	Compute loss . . . . .	4
1.3.4	Backward propagation . . . . .	4
1.4	Gradient Decent . . . . .	5
1.5	Activation Function . . . . .	5
1.6	Pitfalls . . . . .	6
<b>2</b>	<b>Implementation</b>	<b>6</b>
2.1	Object Oriented Programming in MATLAB . . . . .	6
2.2	MATLAB Code . . . . .	6
<b>3</b>	<b>Results</b>	<b>7</b>
3.1	Simple Neural Net . . . . .	7
3.2	Comparison of different hidden layer sizes . . . . .	7
3.3	Multiple hidden layers . . . . .	7
<b>4</b>	<b>Conclusion</b>	<b>7</b>
	<b>Appendix</b>	<b>9</b>

## List of Figures

1	Sample numbers from MNIST [1]. . . . .	2
2	Visualization of a neural network with one hidden layer [2]. . . . .	3
3	A visualization of forward and backward propagation [3]. . . . .	4
4	Visualization of sigmoid and Tanh function . . . . .	5
5	Run times for various neural network architectures [4]. . . . .	6
6	The results from the simple network for the first 1000 epochs. . . . .	7
7	Net accuracy for different layer sizes. . . . .	8
8	The results from the two layer network for the first 2000 epochs. . . . .	8

## Abstract

The purpose of this report is to develop a neural net that can identify handwritten digits in the MNIST database at near human levels of accuracy. The neural net will be developed without the assistance of libraries such as Python's tensor flow or MATLAB's Deep Learning.

The author would like to express his gratitude to Professor Hicken for the suggestion of this project. The author would also like to thank Theodore Ross and Varun Rao for their assistance with artificial neural networks.

## 1 Introduction

**Put a paragraph here** In recent years the Have it solve the MNIST with a simple simple thing then try different layers and stuff

Go over tan h vs sigmoid Explain batch testing

### 1.1 The MNIST database

The Modified National Institute of Standards and Technology database or MNIST database [5] is a database of handwritten numbers used to train image processing systems. It contains 60,000 training images and 10,000 testing images. The database is comprised of images that are made up of a grid of 28x28 pixels. Some of these are seen in figure 1.

A number of attempts have been made to get the lowest possible error rate on this dataset. As of August 2018 the the lowest achieved so far is a error rate of 0.21% or an accuracy of 99.79%. For comparison human can accurately recognize digits at a rate of 98.5% [6].




Figure 1: Sample numbers from MNIST [1].

### 1.2 Artificial neural network

An artificial neural network (referred to as a neural network in this paper) is a computation system that mimics the biological neural networks found in animal brains. A neural network is not an explicit A Neural networks may be trained for tasks, such as the number recognition in this report.

The neural net implemented in this project had an input vector of  $784 \times 1$  and an output vector of  $10 \times 1$ . Different configurations were tried, with one hidden layer of  $250 \times 1$  producing the best results. A visualization of an example neural net is seen in figure 2.




Figure 2: Visualization of a neural network with one hidden layer [2].

### 1.3 Neural Network Walkthrough

The training of a neural network involves four main steps:

1. Initialize weights and biases.
2. Forward propagation
3. Compute the loss
4. Backward propagation

#### 1.3.1 Parameter Initialization

The first step in training a neural net is to initialize the bias vectors and weight matrices. They are initialized with random numbers between 0 and 1, then multiplied by a small scalar around the order of  $10^{-2}$  so that the units are not on the region where the derivative of the activation function are close to zero. The initial parameters should be different values (to keep the gradients from being the same).

There are various forms of initialization such as Xavier initialization or He-et-al Initialization, but a discussion on methods of initialization outside the scope of this paper. In this paper we will stick with random parameter initialization.

#### 1.3.2 Forward Propagation

The next step is the forward propagation. The network takes the inputs from a previous layer, computes their transformation, and applies an activation function. Mathematically the forward propagation at level “i” is represented by equation 1.




Figure 3: A visualization of forward and backward propagation [3].

$$\begin{aligned} z_i &= A_{i-1} * W_i + b \\ A_i &= \phi(z_i) \end{aligned} \quad (1)$$

Where  $z$  is the input vector,  $A$  is the layer,  $W$  is the weights going into the layer,  $b$  the bias, and  $\phi$  the activation function. This process then repeats for the next layer until it reaches the end of the neural net.

### 1.3.3 Compute loss

The loss is simply the difference between the output and the actual value. In this neural net it is computed by equation 2.

$$\text{loss} = A_{i=\text{end}} - y \quad (2)$$

### 1.3.4 Backward propagation

After going forward through the neural net in the forward propagation step, the next step is backwards propagation. Backwards propagation is the updating of the weight parameters via the derivative of the error function with respect to the weights of the neural net. For the output layer this is seen in equation 3 and equation 4 for all other layers.

$$dW_{i=\text{end}} = \phi'(z_{i=\text{end}}) * (A_{i=\text{end}} - y) \quad (3)$$

$$dW_i = \phi'(z_i) * (W_{(i+1)}^T * dW_{(i+1)}) \quad (4)$$

Once these derivatives have been computed, the weights are updated by equation 5

$$W_i = W_i + \alpha * dW_i * A_{(i-1)}^T \quad (5)$$

Where for the first layer  $z_{(i-1)}^T$  will be the input vector and for all the following layers it will be the vector from the previous layer.

Forward and backward propagation are visualized in figure 3




Figure 4: Visualization of sigmoid and Tanh function

## 1.4 Gradient Decent

Also known as steepest decent, gradient decent is a first order optimization algorithm. It is used to find the minimum of a function. Equation 6 shows gradient decent implemented in a neural net.

$$\Delta W(t) = -\alpha \frac{\partial E}{\partial W(t)} \quad (6)$$

Where **put more here!!**

## 1.5 Activation Function

The activation function was previously mentioned as a function used to convert the input signal to the output signal. Activation functions introduce non-linear properties to the neural net's functions, allowing the neural net to represent complex functions [3].

The two most common activation functions used in neural nets for the gradient decent are sigmoid and hyperbolic tangent (Tanh). The formula for Tanh is seen in equation 7, and the formula for it's derivative is seen in equation 8

$$\phi_{\text{Tanh}}(z) = \frac{1 - e^{-2z}}{1 + e^{-2z}} \quad (7)$$

$$\phi'_{\text{Tanh}}(z) = \frac{4}{(e^{-z} + e^z)^2} \quad (8)$$

The formula for the sigmoid function is seen in equation 9, the formula for it's derivative is seen in equation 10.

$$\phi_{\text{Sigmoid}}(z) = \frac{1}{1 + e^{-z}} \quad (9)$$

$$\phi'_{\text{Sigmoid}}(z) = \frac{e^{-z}}{(e^{-z} + 1)^2} \quad (10)$$

The sigmoid and Tanh function are visualized in figure 4.

Both functions have relatively simple mathematical formulas and are differentiable. In this paper the sigmoid function is used over the Tanh function. **WHY?** Sigmoid and Tanh are not the only activation functions. Other functions that should be noted are the Rectified Linear Unit (ReLU) and the Leaky Rectified Linear Unit function. While these functions can perform better than Tanh and Sigmoid, they are more complex and a proper discussion of them is outside the scope of this paper.

**Expalain importance of activation function**

ID	architecture (number of neurons in each layer)	test error for best validation [%]	best test error [%]	simulation time [h]	weights [milions]
1	1000, 500, 10	<b>0.49</b>	0.44	23.4	1.34
2	1500, 1000, 500, 10	<b>0.46</b>	0.40	44.2	3.26
3	2000, 1500, 1000, 500, 10	<b>0.41</b>	0.39	66.7	6.69
4	2500, 2000, 1500, 1000, 500, 10	<b>0.35</b>	0.32	114.5	12.11

Figure 5: Run times for various neural network architectures [4].

## 1.6 Pitfalls

The most important thing to steer clear of is over training. Overtraining occurs when the neural net trains too much to the training data. While it will have a high accuracy for the training data, its performance for the test data will decay, as it has become too well attuned to the training data.

The other problem is the time it takes to train. A three layer neural net can be trained to 97% accuracy within 10 minutes, however it will not improve far beyond that. Larger nets will take longer to train, but will take far longer to train.

# 2 Implementation

## 2.1 Object Oriented Programming in MATLAB

This neural net had to be made without the use of any built in libraries [7] and the code had to be modular [8]. To create code for neural network subject to these constraints the author decided to create their own neural net class in MATLAB.


The MATLAB class `philipNeuralNet.m` was written for this neural net project. It has the parameters `learningRate` and `Level`. The `learningRate` parameter is obviously the learning rate. The `Level` parameter has four parameters attached to it: `W` (weight), `dW` (weight derivative), `z` (input), and `A` (the vector for the layer). By having this class we have avoided hard coding the propagation of the neural net and it is possible to test different neural net architectures on this code.

The MATLAB class also has an activation function and a derivative of the activation function. The `actFunSwitch` variable allows either the Sigmoid or the Tanh function to be selected. Additionally the `enableBias` variable allows for biases to be used or not used in the code's execution.


Finally it has a `outputVector` function that is simply an implementation of the neural net. It takes the input, runs it through the net, and returns the net's output.

## 2.2 MATLAB Code

The MATLAB code first initializes a neural net from given parameters. It obtains the MNIST data from a function [9]. It then uses the `handleTrainNet` function to train the net. This function implements batch training, using the forward and backward propagation functions. It then computes and displays the training error and the testing accuracy after a specified number of runs via the `testAcc.m` function. Once it has done this it plots the accuracy of the neural net via the `plot accuracy` function, generating the plots seen in the results section.



(a) The results for the first 1000 epochs.



(b) The results for the first epochs scaled logarithmically

Figure 6: The results from the simple network for the first 1000 epochs.

## 3 Results

### 3.1 Simple Neural Net

The best results were found for simplex neural net examined; a one hidden layer with 250 nodes a learning rate of 0.1, and no biases. For this simple a test accuracy of 98% was achieved. The first 1000 epochs of this net are visualized in figure 6. To get the 98.37% accuracy it took approximately 12 hours of running the code and over three million neural net evaluations.

### 3.2 Comparison of different hidden layer sizes

From Shure [10], the optimal size for the hidden layer in a three layer neural network is 250 nodes. Comparing the results for hidden layers show in figure 7, different sizes of hidden layer do not have a large effect on the accuracy of these results. However what is different is the time it takes to run each net. The more nodes in a net, the longer it takes for the net to train, as there are more operations to perform. Thus if a neural net with 250 nodes will have the same accuracy as a net with 800 nodes, the first net is preferable, as it will be trained faster.

### 3.3 Multiple hidden layers

For network with two hidden layers of 250 each the best test accuracy was 96.49%. The accuracy from the first 2000 epochs is seen in figure 8.

## 4 Conclusion

The simple neural net achieved an accuracy of 98.37%. This is on par with human recognition, and is about as accurate as a simple neural net can achieve. Higher accuracy rates are achieved via the use of constitutional neural networks, such as the LeNet-5 [?]. Due to the long run time of large multi layered neural networks they were not studied, but could provide a more accurate identification with out convolution.






Figure 7: Net accuracy for different layer sizes.



(a) The results for the first 2000 epochs.



(b) The results for the first epochs scaled logarithmically

Figure 8: The results from the two layer network for the first 2000 epochs.



What is interesting is that the best results occurred with out biases and with one hidden layer. It was expected that adding more complexion to the neural net would increase the accuracy, however this was not the case. As neural nets are very much a trial and error process, it is possible that these more complex nets will achieve a better accuracy with more fiddling.

## References

- [1] Deep learning for classification on the mnist dataset. <https://www.mathworks.com/examples/computer-vision/community/36153-deep-learning-for-classification-on-the-mnist-dataset>, 2018.
- [2] Diagram of an artificial neural network. <https://tex.stackexchange.com/questions/132444/diagram-of-an-artificial-neural-network>, September 2013.
- [3] Daphne Cornelisse. [https://medium.freecodecamp.org/building-a-3-layer-neural-network-from-scratch-fbclid=IwAR1jjh1IsEVdvN0PIeMygzfY2ZG3K-Zata3z\\_jqlfLtQZKKOX6QEbnZABzw](https://medium.freecodecamp.org/building-a-3-layer-neural-network-from-scratch-fbclid=IwAR1jjh1IsEVdvN0PIeMygzfY2ZG3K-Zata3z_jqlfLtQZKKOX6QEbnZABzw), February 2018.
- [4] Luca Maria Gambardella Jurgen Schmidhuber Dan Claudiu Ciresan, Ueli Meier. Deep big simple neural nets excel on handwritten digit recognition. <https://arxiv.org/pdf/1003.0358.pdf>, March 2010.
- [5] Christopher J.C. Burges Yann LeCun, Corinna Cortes. The mnist database of handwritten digits. <http://yann.lecun.com/exdb/mnist/>.
- [6] John Denker Patrice Simard, Yann Le Cun. Efficient pattern recognition using a new transformation distance. <https://papers.nips.cc/paper/656-efficient-pattern-recognition-using-a-new-transformation-distance.pdf>.
- [7] Jason Hicken. Mane 6963 independent study description, December 2018.
- [8] Jason Hicken. Rubric for mane 6963 independent study, December 2018.
- [9] Using the mnist dataset. [http://ufldl.stanford.edu/wiki/index.php/Using\\_the\\_MNIST\\_Dataset](http://ufldl.stanford.edu/wiki/index.php/Using_the_MNIST_Dataset), may 2011.
- [10] Loren Shure. Artificial neural networks for beginners. <https://blogs.mathworks.com/loren/2015/08/04/artificial-neural-networks-for-beginners/>, August 2015.

## Appendix 1 - MATLAB code

### NN\_Master.m

```

1 %% Philip Hoddinott NN
2 % Neural Net for MNIST numbers
3 %% Setup
4 clear all; close all;
5 %% load values
6 % These functions come from
7 % http://ufldl.stanford.edu/wiki/index.php/Using_the_MNIST_Dataset
8 inputValues = loadMNISTImages('train-images.idx3-ubyte');
```

```

9 labels = loadMNISTLabels('train-labels.idx1-ubyte');
10
11 inputValuesTest = loadMNISTImages('t10k-images.idx3-ubyte');
12 labelsTest = loadMNISTLabels('t10k-labels.idx1-ubyte');
13
14 % change labels
15 targetValues = 0.*ones(10, size(labels, 1));
16 for n = 1: size(labels, 1)
17     targetValues(labels(n) + 1, n) = 1;
18 end
19 % traing paramters
20 sizeArr=[250;10];
21 learningRate=.1;
22 batchSize = 100;
23 epochs = 500; numEpoch=1;
24 % net switches
25 enableBias=0; % 0 for off, 1 for on
26 actFunSwitch =0; % 0 for sigmoid, 1 for Tanh
27 % create net
28 pNet=philipNeuralNet(inputValues,sizeArr,learningRate,enableBias,actFunSwitch);
29 % train net
30 [pNet,pNetBest,errorM,testAccM]=handleTrainNet(pNet, inputValues, targetValues, ...
    epochs, batchSize, inputValuesTest, labelsTest,numEpoch,sizeArr);
31 % plot results
32 plotAcc
33
34 function [pNet,pNetBest,errorM,testAccM] = handleTrainNet(pNet, inputValues, ...
    targetValues, epochs, batchSize, inputValuesTest, labelsTest,numEpoch,sizeArr)
35 % handleTrainNet function to train net via batch traning
36 trainingSetSize = size(inputValues, 2); % get traning set size
37 errorM=[]; testAccM=[]; % init values
38
39 n = zeros(batchSize); % init values
40 errorBest=100; accBest=-1; % init values
41
42 figure; hold on; % init figure
43 ylabel('Training Error')
44 xlabel('Epochs')
45 tic
46 for t = 1: numEpoch*epochs
47     for k = 1: batchSize
48         % Select input vector to train on.
49         n(k) = floor(rand(1)*trainingSetSize + 1);
50         % get inputs and targets
51         inputVector = inputValues(:, n(k));
52         targetVector = targetValues(:, n(k));
53         % forward propogation
54         pNet = forwardProp(pNet,sizeArr,inputVector);
55         % backwards Propagation
56         iArr=linspace(length(sizeArr),1,length(sizeArr));
57         pNet=backprop(pNet,sizeArr,inputVector,targetVector);
58     end % end for loop
59
60     % Calculate the error for plotting.
61     error = 0;
62     for k = 1: batchSize
63         inputVector = inputValues(:, n(k));
64         targetVector = targetValues(:, n(k));
65         outputVector = pNet.netOutput(inputVector,sizeArr);

```

```

66         error=error+norm(outputVector- targetVector, 2);
67
68     end
69     error = error/batchSize; errorM=[errorM,error];
70     plot(t,error,'*')
71
72     if error<errorBest
73         pNetBest=pNet; errorBest=error;
74     end
75
76     if mod(t,25)==0 %
77         [numCorrect, numErrors] = testAcc(pNet, inputValuesTest, ...
78             labelsTest,sizeArr);
79         acc=100*(numCorrect) / (numCorrect+numErrors);
80         if acc>accBest
81             accBest=acc;
82         end
83         testAccM=[testAccM,acc];
84         fprintf('Epoch = %d,error = %.4f, best acc = %.4f\n',t, error,accBest)
85         grid on;
86         toc
87     end
88     drawnow % draw error
89 end
90 end
91
92 function pNet = forwardProp(pNet,sizeArr,inputVector)
93     for i=1:length(sizeArr)
94         if i==1
95             pNet.Level(i).z=pNet.Level(i).W*inputVector;
96         else % output
97             pNet.Level(i).z=pNet.Level(i).W* pNet.Level(i-1).A;
98         end
99         pNet.Level(i).A=pNet.actFunc(pNet.Level(i).z+pNet.Level(i).b);%%% NEW
100     end
101 end
102
103 function pNet=backprop(pNet,sizeArr,inputVector,targetVector) % function to ...
104     perform backpropagation
105     learningRate=pNet.learningRate;
106     iArr=linspace(length(sizeArr),1,length(sizeArr));
107     for i=iArr
108         if i==length(sizeArr) % cost at output
109             pNet.Level(i).dW=pNet.dactFunc( pNet.Level(i).z).* ...
110                 (pNet.Level(i).A-targetVector);
111         else % hidden
112             pNet.Level(i).dW = pNet.dactFunc( ...
113                 pNet.Level(i).z.*(pNet.Level(i+1).W'* pNet.Level(i+1).dW);
114             end
115             dz=pNet.dactFunc(pNet.Level(i).z);
116             pNet.Level(i).db=(1/length(dz))* sum(dz,2); %%% NEW
117             %pNet.Level(i).dW=pNet.Level(i).dW*(1/length(pNet.Level(i).dW)); %% NEW
118         end
119     end
120     for i=iArr
121         if i≠1 % output
122             pNet.Level(i).W= ...
123                 pNet.Level(i).W-learningRate*pNet.Level(i).dW*pNet.Level(i-1).A';

```

```

120         else % hidden
121             pNet.Level(i).W= ...
                pNet.Level(i).W-learningRate*pNet.Level(i).dW*inputVector';
122         end
123         if pNet.enableBias==1 % switch for enable bias
124             pNet.Level(i).b=pNet.Level(i).b-learningRate*pNet.Level(i).db; %%% NEW
125         end
126     end
127 end

```

## philipNeuralNet.m

```

1  classdef philipNeuralNet
2      %philipNeuralNet Summary of this class goes here
3      % Detailed explanation goes here
4
5      properties
6          learningRate;
7          Level;
8          enableBias;
9          actFunSwitch;
10     end
11
12     methods
13         function obj = ...
14             philipNeuralNet(inputValues,sizeArr,learningRate,enableBias,actFunSwitch)
15             % initialized values
16             inputDim = size(inputValues, 1);
17
18             for i =1:length(sizeArr)
19                 if i==1
20                     obj.Level(i).W=rand(sizeArr(i),inputDim);
21                     obj.Level(i).W=( obj.Level(i).W)./size( obj.Level(i).W,2);
22                 else
23                     obj.Level(i).W=rand(sizeArr(i),sizeArr(i-1));
24                     obj.Level(i).W=( obj.Level(i).W)./size( obj.Level(i).W,2);
25                 end
26                 obj.Level(i).z=learningRate*rand(sizeArr(i),1);
27                 obj.Level(i).A=learningRate*rand(sizeArr(i),1);
28                 obj.Level(i).dW=learningRate*rand(sizeArr(i),1);
29                 obj.Level(i).b=learningRate*zeros(sizeArr(i),1);
30                 obj.Level(i).db=learningRate*obj.Level(i).b;
31
32             end
33             obj.learningRate=learningRate;
34             obj.enableBias=enableBias;
35             obj.actFunSwitch=actFunSwitch;
36         end
37         function funcVal = actFunc(obj,x)
38             %actFunc activation function
39             % depending on the activation funciton switch, this performs
40             % sigmoid or TanH
41             if obj.actFunSwitch==0 % for sigmoid
42                 funcVal = 1./(1 + exp(-x));
43             elseif obj.actFunSwitch==1 % for tanh
44                 funcVal=tanh(x);

```

```

44         end
45     end
46
47     function funcD = dactFunc(obj,x)
48         %dactFunc activation func derivative
49         % depending on the activation function switch, this performs
50         % derivative of sigmoid or Tanh
51         if obj.actFunSwitch==0
52             funcD = obj.actFunc(x).*(1 - obj.actFunc(x));
53         elseif obj.actFunSwitch==1
54             funcD=(1-tanh(x).^2);
55         end
56     end
57
58     function outputVector = netOutput(pNet, inputVector, sizeArr)
59         % netOutput get the output of the neural net given an input and
60         %
61         % INPUT:
62         % pNet : net
63         % inputVector : input to net
64         % sizeArr : net architecture
65         %
66         % OUTPUT:
67         % outputVector : output of net
68         for i=1:length(sizeArr)
69             if i==1
70                 pNet.Level(i).z=pNet.Level(i).W*inputVector;
71             else
72                 pNet.Level(i).z=pNet.Level(i).W*pNet.Level(i-1).A;
73             end
74             pNet.Level(i).A=pNet.actFunc(pNet.Level(i).z);
75         end
76         outputVector=pNet.Level(i).A;
77     end
78
79 end
80 end

```

## testAcc.m

```

1 function [numCorrect, numErrors] = testAcc(pNet, inputValues, labels, sizeArr)
2     % testAcc test the accuracy of a net using mnist validation set
3     %
4     % INPUT:
5     % pNet : net
6     % inputValues : MNIST Input values for training
7     % labels : MNIST Labels for validation
8     % sizeArr : net architecture
9     %
10    % OUTPUT:
11    % numCorrect : number of correctly classified numbers.
12    % numErrors : number of classification errors.
13    %
14    testSetSize = size(inputValues, 2);
15    numErrors = 0;    numCorrect = 0;
16

```

```
17     for n = 1: testSetSize
18         inputVector = inputValues(:, n);
19         outputVector = pNet.netOutput(inputVector, sizeArr);
20         max = 0; class = 1;
21
22         for i = 1: size(outputVector, 1)
23             if outputVector(i) > max
24                 max = outputVector(i);
25                 class = i;
26             end
27         end
28
29         if class == labels(n) + 1
30             numCorrect = numCorrect + 1;
31         else
32             numErrors = numErrors + 1;
33         end
34     end
35 end
```